

Deep Learning Classification with Tensorflow and Keras

Loren Collingwood

Deep learning is one of the most recent developments in machine learning based off tensorflow python package ported via the keras R package. In essence, we are building off the neural networks approach to machine learning. Overall, this process is most similar to the supervised machine learning algorithms/approach found in RTextTools but is more updated. We will work through an example from the tensorflow/rstudio website. I'll be honest, some of you may not be able to line your computer up to the right specifications in order to get tensorflow and keras to work.

Step 1

Install and download packages, then load example data.

```
options(scipen = 999, digits = 4)

#####
# Packages #
#####

install.packages(c("keras", "purrr", "pins", "tensorflow"))

library(keras)
library(dplyr)

##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
library(ggplot2)
library(purrr)
library(pins)
library(tensorflow)

# Had to create a kaggle account and then download this api key #
# Otherwise you can download the data yourself and store it some place nice
board_register_kaggle(token = "~/Downloads/kaggle.json")

# Pull in the data and movie review .csv file
path <- pins::pin_get("nltkdata/movie-review", "kaggle")[1]
```

```
# Read in the .csv file from the path
df <- readr::read_csv(path)
```

```
## Parsed with column specification:
## cols(
##   fold_id = col_double(),
##   cv_tag = col_character(),
##   html_id = col_double(),
##   sent_id = col_double(),
##   text = col_character(),
##   tag = col_character()
## )
```

```
# take a look at first 6 rows
head(df)
```

```
## # A tibble: 6 x 6
##   fold_id cv_tag html_id sent_id text tag
##   <dbl> <chr>   <dbl>   <dbl> <chr> <chr>
## 1      0 cv000    29590      0 "films adapted from comic books have had~ pos
## 2      0 cv000    29590      1 "for starters , it was created by alan m~ pos
## 3      0 cv000    29590      2 "to say moore and campbell thoroughly re~ pos
## 4      0 cv000    29590      3 "the book ( or \" graphic novel , \" if ~ pos
## 5      0 cv000    29590      4 "in other words , don't dismiss this fil~ pos
## 6      0 cv000    29590      5 "if you can get past the whole comic boo~ pos
```

```
# Look at text 1 for shites and giggles
df$text[1]
```

```
## [1] "films adapted from comic books have had plenty of success , whether they're about superheroes (
```

```
# Count of review variable #
df %>% count(tag)
```

```
## # A tibble: 2 x 2
##   tag      n
##   <chr> <int>
## 1 neg   31783
## 2 pos   32937
```

```
# Take a look at distribution of Words in Each Review
```

```
df$text %>%
  strsplit(" ") %>%
  sapply(length) %>%
  summary()
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      1.0   14.0    21.0   23.1   30.0   179.0
```

Step 2

Create training and testing datasets. We will use these to evaluate model accuracy.

```
# Create Unique id, 80% of dataset; used for separating train/test set
training_id <- sample.int(nrow(df), size = nrow(df)*0.8)
```

```
# Generate Train and Test sets
training <- df[training_id,]
testing <- df[-training_id,]
```

Step 3

Define Text Vectorization Level. This will install miniconda; then install tensorflow(). You might have some problem with installation here depending on your computer setup and your computer knowledge.

```
# Create Text Vectorization Layer #
num_words <- 10000 # you can mix it up here but we'll start with this
max_length <- 50

text_vectorization <- layer_text_vectorization(
  max_tokens = num_words,
  output_sequence_length = max_length,
)

text_vectorization %>%
  adapt(df$text)

# Look at vocabulary/words #
# TODO see https://github.com/tensorflow/tensorflow/pull/34529
head ( get_vocabulary(text_vectorization) )

## [1] "the" "a" "and" "of" "to" "is"

#You can see how the text vectorization layer transforms it's inputs:
text_vectorization(matrix(df$text[1], ncol = 1))

## tf.Tensor(
## [[ 68 2835 30 359 1662 33 91 1056 5 632 631 321 41 7803
## 709 4865 1767 48 7600 1337 398 5161 48 2 1 1808 1800 148
## 17 140 109 90 69 3 359 408 40 30 503 142 0 0
## 0 0 0 0 0 0 0 0 0]], shape=(1, 50), dtype=int64)
```

Step 4

Build models inputs and outputs then send to keras_model() function.

```
# Build Model for Input
input <- layer_input(shape = 1, dtype = "string")

# Output
output <- input %>%
  text_vectorization() %>%
  layer_embedding(input_dim = num_words + 1, output_dim = 16) %>%
  layer_global_average_pooling_1d() %>%
  layer_dense(units = 16, activation = "relu") %>%
  layer_dropout(0.5) %>%
  layer_dense(units = 1, activation = "sigmoid")
```

```

model <- keras_model(input, output)

# Hidden Units and Loss Optimizer #
model %>% compile(
  optimizer = 'adam',
  loss = 'binary_crossentropy',
  metrics = list('accuracy')
)

```

Step 5

Train the model real good using the fit() function.

```

#####
# Train the Model #
#####

history <- model %>% fit(
  training$text,
  as.numeric(training$tag == "pos"),
  epochs = 10,
  batch_size = 512,
  validation_split = 0.2,
  verbose=2
)

```

Step 6

Then look at the results and extract predictions on virgin texts as needed.

```

results <- model %>% evaluate(testing$text, as.numeric(testing$tag == "pos"), verbose = 0)
results

```

```

##      loss accuracy
## 0.6015 0.6780

```

```


# Look at the within model accuracy/loss
plot(history)

```

```

## `geom_smooth()` using formula 'y ~ x'

```



```
#####
# Predict Classes #
#####

# This is what you would do onto virgin text

head(pred_out <- as.data.frame(predict(model,
  x = testing$text)))

##      V1
## 1 0.2534
## 2 0.4921
## 3 0.3944
## 4 0.4229
## 5 0.4727
## 6 0.2452

# Create the dummy/binary where 1 = pos, 0 = negative (rating)
pred_out$pred_bin <- ifelse(pred_out$V1 >= 0.5, 1, 0)

# Just add on the original truth #
pred_out$truth <- as.numeric(testing$tag == "pos")

# Look at Confusion Matrix #
table(pred_out$pred_bin, pred_out$truth)

##
##      0      1
```

```
##    0 4457 2259
##    1 1909 4319
# Run the numbers real nice #
(4371 + 4404) / nrow(testing)
## [1] 0.6779
```